Almond classification: Deep Learning Algorithms

Leonardo Aldecocea Colomar

August 26, 2024

Executive Summary:

This study compared three deep learning models (MLP, CNN, and RNN-LSTM) for a three-class classification task involving MAMRA, REGULAR, and SANORA categories. The RNN-LSTM model demonstrated superior performance with an overall accuracy of 80%, significantly outperforming both the MLP (33% accuracy) and CNN (76% accuracy) models. The RNN-LSTM model showed particularly strong results in classifying the MAMRA category, with a 91% F1-score. These findings suggest that the RNN-LSTM architecture is the most suitable for this specific classification task.

Introduction:

In the field of machine learning and artificial intelligence, choosing the right model architecture for a given classification task is crucial. This study aims to compare the performance of three popular deep learning architectures - Multilayer Perceptron (MLP), Convolutional Neural Network (CNN), and Recurrent Neural Network with Long Short-Term Memory (RNN-LSTM) - on a specific three-class classification problem. The classes in question are MAMRA, REGULAR, and SANORA, though the exact nature of these categories is not specified in the provided data.

The objective of this comparison is to determine which model architecture is most effective for this particular classification task, providing insights that can guide future model selection and development in similar scenarios.

Methodology:

Three deep learning models were implemented and trained on the same dataset:

1. Multilayer Perceptron (MLP): A feedforward neural network with multiple layers of perceptrons.

2. Convolutional Neural Network (CNN): A deep learning architecture typically used for processing grid-like data, such as images.

3. Recurrent Neural Network with Long Short-Term Memory (RNN-LSTM): A type of neural network designed to recognize patterns in sequences of data.

Each model was trained on a dataset containing instances of MAMRA, REGULAR, and SANORA classes. The dataset was split into training and testing sets, with the test set containing 2,803 instances (MAMRA: 933, REGULAR: 927, SANORA: 943).

After training, each model's performance was evaluated using standard classification metrics: precision, recall, and F1-score for each class, as well as overall accuracy. These metrics were calculated based on the model's predictions on the test set.

Results:

1. MLP Model:
   - Overall Accuracy: 33%
   - Poor performance across all classes, with 0% F1-score for MAMRA and SANORA
   - Overfitting to the REGULAR class
2. CNN Model:
   - Overall Accuracy: 76%
   - MAMRA: F1-score 0.88
   - REGULAR: F1-score 0.72
   - SANORA: F1-score 0.67
3. RNN-LSTM Model:
   - Overall Accuracy: 80%
   - MAMRA: F1-score 0.91
   - REGULAR: F1-score 0.75
   - SANORA: F1-score 0.76

The RNN-LSTM model demonstrated the best overall performance, with high accuracy and balanced performance across all three classes.

Conclusion/Inferences:

1. Model Performance: The RNN-LSTM model significantly outperformed both the MLP and CNN models, suggesting that this architecture is most suitable for the given classification task. The sequential nature of LSTM networks may be particularly well-suited to capturing relevant patterns in the data.
2. Class-specific Performance: All models performed best on the MAMRA class, indicating that this class may have the most distinctive features. The REGULAR and SANORA classes proved more challenging to classify, suggesting potential similarities between these classes that the models found difficult to distinguish.

3. MLP Limitations: The poor performance of the MLP model highlights the importance of choosing an appropriate architecture for the specific data and task at hand. The MLP's inability to capture complex patterns in the data resulted in severe overfitting to a single class.

4. Potential for Improvement: While the RNN-LSTM model performed well, there's still room for improvement, particularly in classifying REGULAR and SANORA instances. Further fine-tuning of the model architecture, hyperparameter optimization, or the use of ensemble methods could potentially enhance performance.

5. Future Directions: Given the strong performance of the RNN-LSTM model, future work could focus on optimizing this architecture further. Additionally, investigating the specific features that contribute to the model's decisions could provide valuable insights into the nature of the classification task and the distinctions between classes.

In conclusion, this study demonstrates the superiority of the RNN-LSTM architecture for this specific three-class classification task, providing a strong foundation for future model development and application in similar contexts.